

DFA memory-bound face: CUDA/Warp drop past L2; Triton flat (model-bound)

