

DFA memory-bound face: CUDA/Warp peak at L2 then drop; Triton flat

