

From Diagnosis to Cure: Decomposing the Tile-SPMD Abstraction Regret on Irregular Automata

Anonymous

Abstract—Tile-based GPU DSLs (Triton) trail hand-written CUDA by $\sim 10\times$ on irregular finite-automata workloads. Prior work attributed this “abstraction regret” to the execution paradigm (tile/SPMD vs. thread-SIMT) rather than abstraction height. We turn that diagnosis into an anatomy and a cure: we decompose the regret into three independently measured components and identify, to the instruction level, the single missing primitive. On a work-efficient NFA worklist the $10\times$ anchor decomposes as (A) a launch-configuration artifact ($\sim 3.4\times$, default `num_warps`), (B) a lane-packing-recoverable component (warp-redundant scalar execution), and (C) an irreducible $\sim 2\times$ residual. Component C is *not* instruction count, occupancy, or memory bandwidth: at matched occupancy and fewer warp-instructions than CUDA, and sitting below both the issue and bandwidth roofline ceilings, a per-lane Triton worklist still spends $15.3\times$ more cycles in the dependent-load stall and issues at 9.9% vs. 41% of peak—it is latency-bound. The root cause is abstraction-denied intra-warp memory-level parallelism: a CUDA warp’s lanes issue independent in-flight loads that overlap, whereas a lock-step tile serializes the dependent next-state load. We show the residual is *regime-dependent*: for the memory-bound DFA it closes (lane-packed Triton matches CUDA, $1.05\times$, once the table spills to DRAM). We specify the missing IR primitive—a `scalar_program` region that lowers each lane to an independent instruction stream—and *implement its lowering*: the same idiomatic per-lane source, lowered to the thread model, runs $4.2\times$ faster than the tile lowering and matches or exceeds hand-written CUDA, closing the residual by construction. Every kernel is correctness-gated against a CPU oracle; every number traces to a versioned CSV.

I. INTRODUCTION AND CONTRIBUTIONS

A companion study [1] shows that on irregular automata the Triton→CUDA gap follows the execution paradigm, not the abstraction height. It leaves open *why*, mechanistically, and *what* would close it. This paper answers both. Our contributions:

- 1) A **falsifiable decomposition** of the tile-SPMD automata regret into artifact, recoverable, and irreducible, each measured and Nsight-attributed (Sec. III, Fig. 1).
- 2) Identification of the irreducible residual as **abstraction-denied intra-warp latency hiding**—not redundancy, occupancy, instruction count, or bandwidth—*measured*: fewer warp-instructions than CUDA, same occupancy, below both roofline ceilings, yet $15.3\times$ the dependent-load stall and 9.9% vs. 41% issue activity (Sec. III-D).
- 3) A demonstration that the residual is **regime-dependent** (latency-bound NFA vs. memory-bound DFA), unifying both faces through one latency mechanism (Sec. IV).

- 4) A **launch-configuration finding** (default `num_warps=4` inflates the worklist regret $\sim 3.4\times$), a methodological caution for DSL benchmarking (Sec. III-B).
- 5) The **missing-primitive design** + a bound on its payoff + the thread-model existence proof, distinguished from Gluon’s per-thread *layout* (Sec. V).
- 6) **The cure, implemented**: we lower the same idiomatic per-lane source to the thread model and measure that it runs $4.2\times$ faster than the tile lowering and matches/exceeds hand-CUDA, closing the residual by construction (Sec. V-A, Fig. 6).

II. BACKGROUND AND METHOD

NFA simulation (a work-efficient worklist iterating the active set via `ffs` over a bitmask) is control-flow/latency-bound; DFA simulation (one dependent table gather `trans[cur*256+sym]` per symbol) is memory-bound. CUDA and NVIDIA Warp [2] are thread-SIMT (one thread per string); Triton [3] is tile/SPMD (a program operates on tiles). Correctness gates speed: every kernel/config is validated bit-for-bit against a CPU oracle before any throughput is reported. We report medians over warmup+9 reps at a GPU-saturating batch [4] (timings on GPUs are non-Gaussian, so we avoid $\text{mean}\pm\text{std}$ and best-of- N), and sweep batch because the lane-packing benefit is occupancy-gated. Every mechanism claim is backed by Nsight Compute, not just wall-clock: the issue-stall composition (the `long_scoreboard` reason isolates dependent-load waits), achieved occupancy, issue activity, warp- and thread-instruction counts, and DRAM/L2 traffic. We report negative and partial results and correct two of our own intermediate hypotheses (Secs. III, III-D). Hardware: RTX 4070 (sm_89, 46 SMs), Triton 3.5.1, torch 2.9.1+cu128, CUDA 13.3 / driver 580.

III. THE DECOMPOSITION (NFA WORKLIST)

A. The anchor

A work-efficient register-resident worklist (≤ 64 states, batch 4096, 12 configs): Triton 22 Gbps vs. CUDA 227 Gbps = $10.1\times$ **median** ($9.4\text{--}12.3\times$).

B. Component A: launch-configuration artifact

The Triton worklist defaults to `num_warps=4`: four warps per program redundantly process the *same* string. Sweeping `num_warps` $\in \{1, 2, 4, 8\}$: `nw=1` vs. `nw=4` = $2.77\times$ @4096

→ 3.44×@16384 → 3.69×@65536 (each doubling ~halves throughput). This inflates prior worklist regret numbers and must be disclosed.

C. Component B: warp redundancy, recoverable by lane-packing

The scalar Triton program executes warp-uniformly: `thread_inst_per_inst` = 32.00 (Triton) vs. 30.34 (CUDA), the same number with opposite meaning—CUDA’s 32 lanes run 32 *different* strings, Triton’s run the *same* one—yielding ~90× more warp-instructions. Lane-packing (32 strings into the 32 lanes, exploiting the shared CSR so inner-loop bounds stay scalar) removes it. On the dense scan, pure lane-packing recovers 3.2× → 9.8× → 19.4× (batch 4096/16384/65536) toward the ideal 32×; it is **occupancy-gated** (Fig. 2). *Correction*: we first framed the 90× warp-redundancy as the bottleneck; Nsight shows it is removed ~26× yet throughput moves only 3.8×—it was largely hidden by occupancy.

D. Component C: the irreducible residual

A per-lane Triton worklist (each lane its own `ffs` + per-lane CSR gather, no cross-lane reduce, no active-set union) beats the union variant 1.27× but is still 0.51× of CUDA (≈ 2× slower). Nsight is decisive (Fig. 3): it has *fewer* warp-instructions than CUDA (4.66M vs. 5.07M) and the same occupancy (22.5%), yet is 3.3× slower because issue activity is 9.9%. Raising occupancy does not help—a BLOCK sweep {32, 64, 128, 256} only *worsens* it (0.49× → 0.31×): bigger lock-step tiles add divergence. **Direct measurement** confirms the mechanism: at matched occupancy the per-lane Triton kernel spends 15.3× more cycles in the `long_scoreboard` stall (waiting on a dependent memory load) than CUDA, and the issue-rate ratio (4.1×) tracks the throughput ratio (3.5×), consistent with Little’s law [5], [6]. The instruction roofline (Fig. 4) [7] settles “did you write a worse kernel?”: both kernels sit far below the issue *and* bandwidth ceilings (8.3%/27% vs. 32%/3.4% of peak), and Triton issues fewer warp-instructions yet runs 3.5× slower—it is structurally latency-bound. *Honest refinement*: the per-lane kernel does not issue *fewer* memory requests—it issues 26–84× *more* (masked inactive-lane gathers)—so the tile model pays twice: excess masked traffic *and* no dependent-load hiding. Root cause: a CUDA warp’s lanes are independent in-flight-load generators (memory-level parallelism); a lock-step tile serializes the dependent next-state load, so latency hides only across warps (occupancy-bound). This is distinct from branch/memory divergence [8]: the lanes can be perfectly coalesced and still stall on a dependent load. We distinguish it from the *hardware* fix of intra-warp stalls [9]: the same SM is fast under CUDA and slow under Triton at equal occupancy and fewer instructions, so the loss is abstraction-imposed.

Table I summarizes the multiplicative decomposition: removing the launch artifact and the warp redundancy each recovers 3.7×, leaving the irreducible 1.9× latency residual. By Little’s law (concurrency = throughput × latency), the

TABLE I
THE WORKLIST REGRET DECOMPOSES MULTIPLICATIVELY (BATCH 16384, ≤ 64 STATES, ORACLE-GATED). EACH STAGE REMOVES ONE COMPONENT; 28 → 735 IS $3.7 \times 3.7 \times 1.9 = 26\times$.

Kernel	Gbps	Component removed
Triton worklist (nw=4)	28	—
+ nw=1	104	launch artifact (3.7×)
+ lane-packing	383	warp redundancy (3.7×)
CUDA worklist (thread-SIMT)	735	irreducible latency (1.9×)

latency-bound kernel’s throughput is set by the in-flight-request concurrency it can sustain; the per-lane Triton kernel’s 9.9% issue activity and 15.3× dependent-load stall mean it sustains far less effective concurrency than CUDA at the same occupancy, and the issue-rate ratio (4.1×) tracking the throughput ratio (3.5×) is exactly the relation Little’s law predicts when latency is fixed.

IV. REGIME-DEPENDENCE (DFA): THE UNIFICATION

The DFA (memory-bound) decomposes the same way but the residual *closes in the DRAM regime* (Fig. 5). Across table sizes (cache→L2→DRAM): scalar Triton DFA is flat ~29 Gbps (independent confirmation of the scalar-gather ceiling); lane-packing gives ~12× over it; and the lane-packed/CUDA ratio is 0.55–0.62× in cache but 1.05× at a 16 MB (>L2) table—lane-packed Triton *matches* CUDA. Mechanism (DRAM utilization only 14.6%, so memory-latency not saturated bandwidth): at moderate (cache) latency CUDA’s intra-warp hiding wins ~1.7×; at DRAM latency both must hide latency via cross-warp parallelism and *converge*. So the NFA’s fundamental ~2× residual and the DFA’s closeable gap are one mechanism at different latency scales.

V. THE CURE: THE MISSING PRIMITIVE

Design. A `scalar_program` region (or per-lane `serial_range`) in the tile DSL that lowers the marked code to a per-thread independent instruction stream, giving each lane independent control flow and intra-warp latency hiding. The diagnosis names what it must provide: independent per-lane `ffs/while`, per-lane data-dependent loop bounds, register-resident per-lane bitset, scalar gather. **Bound.** It closes the latency-bound residual (up to ~2× on the NFA) and nothing extra in the already-converged DRAM-DFA regime—a regime-dependent, falsifiable prediction. **Existence proof.** The thread model realizes it: CUDA achieves 1.0× and the high-level Warp DSL [2] reaches 0.9×. **Not already in Gluon.** A runnable probe shows Triton’s low-level Gluon frontend exposes per-thread *layout* [10], not per-lane *control flow*: `gl.load` returns a layout-typed block, never a scalar, so the data-dependent CSR loop cannot even be lowered (compile error captured verbatim). Layout control ≠ control-flow divergence.

TABLE II
THE SAME MISSING PER-LANE PRIMITIVE ACROSS IRREGULAR
WORKLOADS (POSITIONING).

Workload	Per-lane irregularity	Missing capability
NFA/DFA automata	active-set <code>ffs</code> ; data-dep. CSR loop	per-lane <code>while</code> + scalar gather
Var-length / entropy decode	data-dep. bit consumption	per-lane loop + bit-cursor
Ragged batch / MoE routing	variable seq. length; routing	per-lane bounds + scatter
Sparse frontier / BFS [11]	variable vertex degree	per-lane neighbor loop + scatter
Tokenization / lexing	data-dep. longest-match	per-lane DFA scan

A. The cure, implemented

We do not stop at proposing the primitive: we implement its lowering and measure that it closes the residual. A function `lower_scalar_program_to_cuda` takes the *same* idiomatic per-lane body the Triton tile kernel expresses (active-set `ffs` worklist, data-dependent CSR loop, scalar gather, register bitset) and lowers it to a per-thread CUDA kernel (one thread = one string), compiled with `nvcc`. Oracle-validated bit-for-bit; throughput on no-accept automata (sustained full-length scan, removing the early-exit confound), ≤ 64 states, batch 16384. The same per-lane program lowered to *threads* (SP) vs. *tiles* (Triton WP2) differs by $4.2\times$ (median; SP ≈ 1555 vs. WP2 ≈ 378 Gbps), robust across batch, and SP *matches or exceeds* hand-written CUDA (SP/CU = $2.15\times$, since the generated kernel is a minimal thread worklist). So the residual is not merely diagnosed but *closed by construction*: supplying the per-lane lowering recovers—and surpasses—thread-model performance from the identical front-end source. (Integrating this lowering into Triton’s MLIR is engineering; the lowering’s effect is what we validate.)

B. Generality (positioning)

The same per-lane capability recurs across irregular GPU workloads (Table II); each needs per-lane data-dependent control flow with scatter/gather that lock-step tiles gate the same way. The contribution is a *capability* result; we keep the empirical depth on automata.

VI. METHODOLOGICAL INTEGRITY

Our analysis corrected two of its own intermediate hypotheses, which we report rather than hide. First, an early reading framed the $\sim 90\times$ warp-instruction redundancy (Sec. III) as the bottleneck; profiling the lane-packed kernel showed the redundancy is removed $\sim 26\times$ yet throughput moves only $3.8\times$ —it was largely hidden by occupancy, so warp-redundancy is real but not load-bearing. Second, we predicted (following the latency-hiding literature) that the per-lane Triton kernel would sustain *fewer* in-flight memory requests than CUDA; direct measurement refuted this—it issues $26\text{--}84\times$ *more* (masked inactive-lane gathers) and still stalls, so the residual is the inability to *hide* dependent-load latency, not a request-count

deficit. We therefore phrase the mechanism via stall composition and issue activity, not request counts. Both corrections sharpened, rather than weakened, the central claim.

VII. THREATS TO VALIDITY

Single GPU. Results are on one RTX 4070; the *mechanism* (intra-warp latency hiding, the issue-activity differential, the DRAM convergence) is a property of the SIMT execution model and the tile-vs-thread lowering, hence architecture-general, but the absolute decomposition factors should be re-measured on an A100/H100 (larger L2 shifts the DFA crossover). *Prototype scope.* A multi-word generalization of the lane-packed kernel (NWORDS int64 words/lane) is oracle-validated to 256 states, so the prototype is not artificially limited; but the clean $\sim 2\times$ residual is a *register-resident-regime* (≤ 64 -state) result—past 64 states the CUDA register worklist itself degrades (register pressure), so the head-to-head is confounded by both baselines losing the register-resident advantage, and ANMLZoo-scale automata (Brill 42661 = 667 words) make the per-lane multi-word scatter explode—itsself a manifestation of the tile limitation we study. Real automata run oracle-valid but sub-Gbps (an *algorithmic* cost, orthogonal to the regret). *Tuning.* `num_warps` is a knob; we disclose the full sweep rather than a single config, rebutting the autotuning counter-thesis [12]: the Triton/Gluon same-MLIR control plus the disclosed sweep separate *tuning* (which closes component A and part of B) from *expressibility* (the residual C), with a probe that is falsifiable by construction.

VIII. REPRODUCIBILITY

Correctness gates every measurement: each kernel/config is checked bit-for-bit against the CPU oracle (NFA `reference.py` / `simulate_dfa`) before any throughput is reported. Every number in this paper traces to a versioned CSV; all figures regenerate from those CSVs by a single script, so the plots cannot drift from the measurements. The Gluon expressibility claim is a runnable probe that exits 0 on the expected compile failure and 1 if a future Gluon ever compiles it—falsifiable by construction. The artifact (kernels, oracle, sweeps, profiling commands, CSVs, and the `scalar_program` lowering) supports one-command regeneration; all figures here are emitted from the versioned CSVs by a single script.

IX. RELATED WORK

GPU automata form an *algorithmic* axis, all single-implementation CUDA: ngAP’s non-blocking multi-symbol processing [13], HybridSA’s bit-parallel shift-and [14], Bit-Gen’s Parabix bitstreams [15], AsyncAP’s input-symbol asynchrony [16], AutomataBLAS’s AP-as-SpMV [17], and the Hopps sparsity ASIC [18]. Each *varies the algorithm* on one CUDA implementation; none compares GPU programming models or holds the algorithm fixed across DSLs. The sole IR-for-automata work [19] lowers regex to one fixed domain-specific accelerator—a hardware target, not a GPU-paradigm/expressibility study, and with no tile-vs-thread cost.




Fig. 1. NFA worklist throughput (batch 16384, ≤ 64 states) climbs as each regret component is removed: $nw=4 \rightarrow nw=1$ (launch artifact), $\rightarrow$ lane-packing (warp redundancy), $\rightarrow$ CUDA (irreducible latency). All oracle-gated.

Our axis is orthogonal: fixed algorithm, varied programming model. **Tile GPU DSLs:** Triton [3] and the 2024–26 wave exposing more thread/warp control. The closest, Gluon’s Linear Layouts [10], give per-thread *data placement* but not per-lane *control flow* (our probe shows the data-dependent loop cannot lower); Descend [20] makes the thread hierarchy first-class but offers no tile drop-to-scalar escape. Warp [2] is the thread-SIMT existence proof. **Mechanism:** we ground the residual in latency-hiding / MLP-vs-occupancy analysis [5], [6] and place both kernels on the instruction roofline [7]; we distinguish it from divergence [8] and from the *hardware* intra-warp-stall fix of Subwarp Interleaving [9] (we attribute the same stall to the DSL, at equal occupancy and fewer instructions). **Methodology and lineage:** rigorous benchmarking [4], roofline [21], and the performance-portability metric [22] (per-hardware; we measure per-*capability* regret); we rebut the autotuning counter-thesis [12] via the Triton/Gluon same-MLIR control and the disclosed `num_warps` sweep. The gap we fill: no constructive, instruction-level account of *why* a tile DSL pays on irregular control flow, naming the missing primitive, with a regime-dependent residual.

X. CONCLUSION

The abstraction regret on irregular automata is not a monolith: a launch artifact, a recoverable redundancy, and an irreducible intra-warp-latency-hiding residual whose size is set by the latency regime. The cure is a named, bounded IR primitive whose lowering we implement and measure—it closes the residual by construction (4.2 $\times$ over the tile lowering); integrating it into Triton’s MLIR is the landmark future step.

REFERENCES

- [1] Anonymous. The two faces of abstraction regret: Execution paradigm over abstraction height on irregular automata, 2026. Companion diagnosis paper.
- [2] NVIDIA. Warp: A python framework for high-performance GPU simulation and graphics, 2022. <https://github.com/NVIDIA/warp>.




Fig. 2. Lane-packing’s benefit is occupancy-gated: the pure-packing speedup grows 3.2 $\times \rightarrow$ 19.4 $\times$ with batch (toward the ideal 32 $\times$) as more warps fill the GPU.

Same occupancy & warps, fewer warp-inst — yet 3.3 $\times$ slower (latency-bour




Fig. 3. Per-lane Triton vs. CUDA (16384 strings, 32 states): same occupancy, fewer warp-instructions, yet 3.3 $\times$ slower at 9.9% vs. 41% issue activity—latency-bound.

- [3] Philippe Tillet, H. T. Kung, and David Cox. Triton: An intermediate language and compiler for tiled neural network computations. In *MAPL*, 2019.
- [4] Torsten Hoefer and Roberto Belli. Scientific benchmarking of parallel computing systems. In *SC*, 2015.
- [5] Vasily Volkov. *Understanding Latency Hiding on GPUs*. PhD thesis, University of California, Berkeley, 2016. Tech. Rep. UCB/EECS-2016-143.
- [6] Sunpyo Hong and Hyesoon Kim. An analytical model for a GPU architecture with memory-level and thread-level parallelism awareness. In *ISCA*, 2009.
- [7] Nan Ding and Samuel Williams. An instruction roofline model for GPUs. *Concurrency and Computation: Practice and Experience*, 2022.
- [8] Wilson W. L. Fung, Ivan Sham, George Yuan, and Tor M. Aamodt. Dynamic warp formation and scheduling for efficient GPU control flow. In *MICRO*, 2007.
- [9] Sana Damani et al. GPU subwarp interleaving. In *HPCA*, 2022.
- [10] Triton contributors. Linear layouts: Robust code generation of efficient tensor computation using $\mathbb{F}_2$, 2025. arXiv:2505.23819.
- [11] Yangzihao Wang, Andrew Davidson, Yuechao Pan, Yuduo Wu, Andy Riffel, and John D. Owens. Gunrock: A high-performance graph processing library on the GPU. In *PPoPP*, 2016.
- [12] Burkhard Ringlein, Thomas Parnell, and Radu Stoica. GPU performance portability needs autotuning, 2025. arXiv:2505.03780.
- [13] Tianao Ge, Tong Zhang, and Hongyuan Liu. ngap: Non-blocking large-scale automata processing on GPUs. In *ASPLOS*, 2024.




Fig. 4. Instruction roofline (16384 strings, 32 states): both kernels sit far below the issue *and* DRAM-bandwidth ceilings (8.3%/27% for Triton, 32%/3.4% for CUDA), so neither is instruction- nor bandwidth-bound—the residual is latency.




Fig. 5. DFA: lane-packed Triton trails CUDA $\sim 1.7\times$ while the table is cache-resident but *matches* it ($1.05\times$) once the 16 MB table spills past L2—both then rely on cross-warp parallelism. The residual is regime-dependent.




Fig. 6. The cure, implemented (batch 16384, ≤ 64 states): the *same* per-lane source lowered to threads (`scalar_program`) runs 4.2 $\times$ the tile lowering (Triton WP2) and matches/exceeds hand-CUDA—the residual closed by construction.

- [14] HybridSA authors. HybridSA: GPU acceleration of multi-pattern regex matching using bit parallelism. *Proc. ACM Program. Lang. (OOPSLA)*, 2024.
- [15] Tianao Ge, Hongyu Chu, and Hongyuan Liu. Interleaved bitstream execution for multi-pattern regex matching on GPUs. In *MICRO*, 2025.
- [16] Hongyuan Liu, Sreepathi Pai, and Adwait Jog. Asynchronous automata processing on GPUs. *Proc. ACM Meas. Anal. Comput. Syst. (POMACS)*, 2023.
- [17] AutomataBLAS authors. Advancing matrix operations for high-performance and memory-efficient automata processing on GPUs. *ACM Trans. Archit. Code Optim. (TACO)*, 2025.
- [18] Yufeng Du, Joel S. Emer, and Daniel Sánchez. Hopps: Leveraging sparsity to accelerate automata processing. In *ASPLOS*, 2025.
- [19] Filippo Somaini, Luca Carloni, Giovanni Agosta, Marco D. Santambrogio, and Davide Conficconi. Combining MLIR dialects with domain-specific architecture for efficient regular expression matching. In *CGO*, 2025.
- [20] Bastian Köpcke, Sergei Gorlatch, and Michel Steuwer. Descend: A safe GPU systems programming language. In *PLDI*, 2024.
- [21] Samuel Williams, Andrew Waterman, and David Patterson. Roofline: An insightful visual performance model for multicore architectures. *Commun. ACM*, 2009.
- [22] S. J. Pennycook, J. D. Sewall, and V. W. Lee. A metric for performance portability, 2016. arXiv:1611.07409.